

Hybrid Logical Clocks

Sandeep Kulkarni, Murat Demirbas, Deepak
Madeppa, Bharadwaj Avva, Marcelo Leone

Vaibhav Bhagee

Logical Clocks and Vector Clocks

- Disregards the physical notion of time
- Progresses with the occurrence of events and not passage of time
- Captures the happens before relationship

Physical Clocks

- Leverage the system clocks at the nodes
- Issues of clock drift and clock jumps
- Periodic requirement of clock synchronisation
- Uncertainty intervals leading to inability to order events

True Time

- Proposed by google with Spanner
- Tight clock synchronisation using GPS clocks and Atomic Clocks
- Returns intervals [earliest, latest] instead of exact time
- Introduces delays if intervals overlap

Hybrid Logical Clocks

- Refines Physical and Logical Clock
- Preserves the causal ordering property of logical clocks

$$e \text{ hb } f \Rightarrow l.e < l.f,$$

- HLC timestamps are close to the NTP clock

$$|l.e - pt.e| \text{ is bounded.}$$

- Bounded space requirement
- Does not have clock jumps leading to non monotonic time updates

HLC Algorithm

Initially $l.j := 0; c.j := 0$

Send or local event

$l'.j := l.j;$

$l.j := \max(l'.j, pt.j);$

If $(l.j = l'.j)$ then $c.j := c.j + 1$

Else $c.j := 0;$

Timestamp with $l.j, c.j$

Receive event of message m

$l'.j := l.j;$

$l.j := \max(l'.j, l.m, pt.j);$

If $(l.j = l'.j = l.m)$ then $c.j := \max(c.j, c.m) + 1$

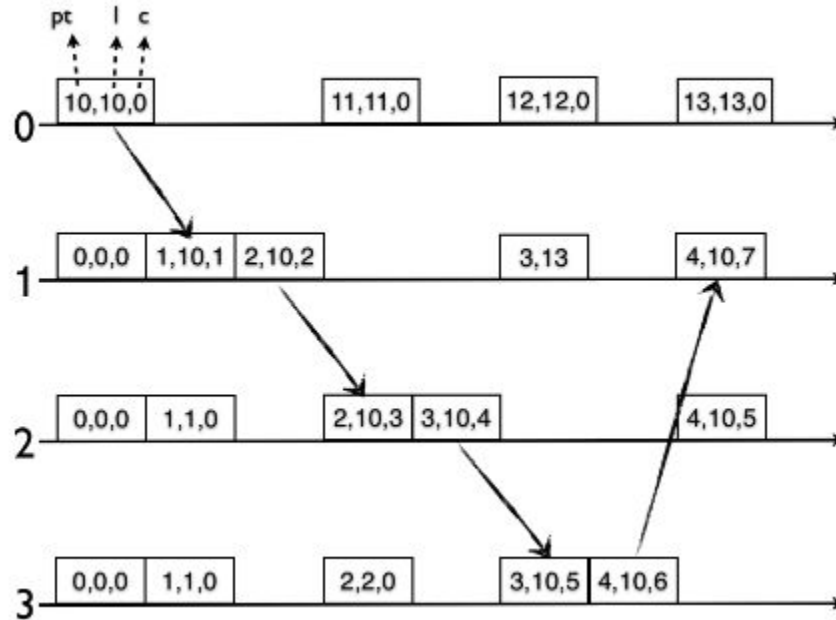
Elseif $(l.j = l'.j)$ then $c.j := c.j + 1$

Elseif $(l.j = l.m)$ then $c.j := c.m + 1$

Else $c.j := 0$

Timestamp with $l.j, c.j$

HLC Algorithm Example



HLC Properties

Theorem 1. *For any two events e and f , $e \text{ hb } f \Rightarrow (l.e, c.e) < (l.f, c.f)$* \square

Theorem 2. *For any event f , $l.f \geq pt.f$* \square

HLC Properties

Theorem 3. *$l.f$ denotes the maximum clock value that f is aware of. In other words,*

$$l.f > pt.f \Rightarrow (\exists g : g \text{ hb } f \wedge pt.g = l.f)$$

- Proof by induction
- Case when f is a send event and e is the previous event
- Case when f is a receive event

$$l.e > pt.e \Rightarrow (\exists g : g \text{ hb } e \wedge pt.g = l.e)$$

- e is the previous event on the same node and m is the received message

$$l.f > pt.f \Rightarrow (\exists g : g \text{ hb } f \wedge pt.g = l.f).$$

HLC Properties

For any event f , $|l.f - pt.f| \leq \epsilon$

- Proof using clock synchronisation constraints
- We cannot have events e and f such that

$$e \text{ hb } f \text{ and } pt.e > pt.f + \epsilon$$

- By the previous theorem

$$l.f > pt.f \Rightarrow (\exists g : g \text{ hb } f \wedge pt.g = l.f).$$

HLC Properties

Theorem 4. For any event f ,

$$c.f = k \wedge k > 0$$

$$\Rightarrow (\exists g_1, g_2, \dots, g_k : \\ (\forall j : 1 \leq j < k : g_j \text{ hb } g_{j+1}) \\ \wedge (\forall j : 1 \leq j \leq k : l.(g_j) = l.f) \\ \wedge g_k \text{ hb } f)$$

- Proof by induction
- Cases: f being a send or a receive event
- Idea is similar to the property of logical clocks
- Number of events causally preceding the current event with same **pt**

HLC Properties

Corollary 2. *For any event f ,*
$$c.f \leq |\{g : g \text{ hb } f \wedge l.g = l.f\})|.$$

Corollary 3. *For any event f , $c.f \leq N * (\epsilon + 1)$*

- First corollary directly follows from the proof
- For proof of second, we use the fact that due to clock synchronisation

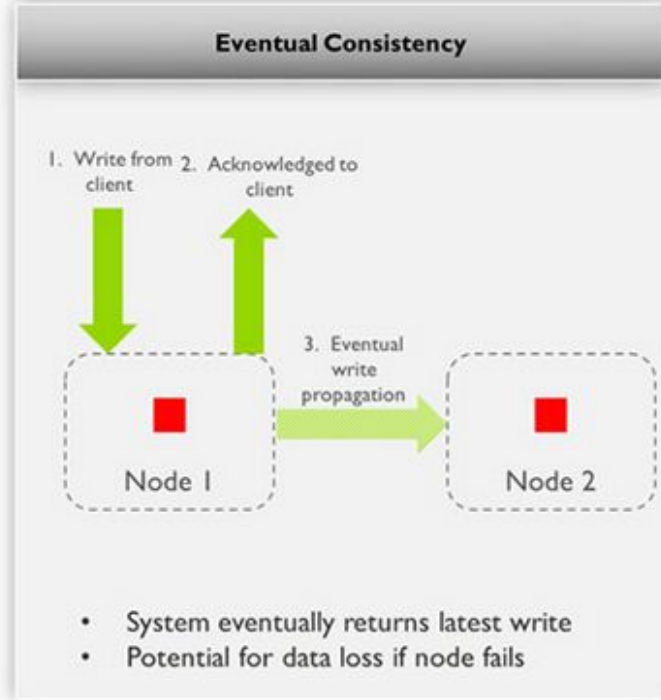
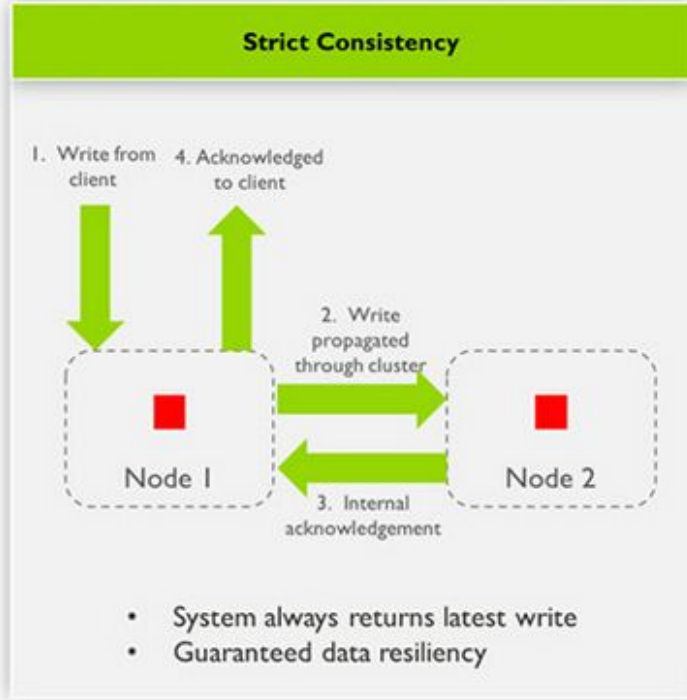
$$[l.f, l.f + \epsilon]$$

- Physical time gets incremented by at least 1 between two events on a node

Applications of HLC: Database Consistency

- **Atomicity:** Transactions are a single unit, either going to completion or not.
- **Consistency:** Database should go from one consistent state to another.
- **Isolation:** Concurrent transaction execution mapped to a serial order
- **Durability:** Once a commit happens, the changes should be persistent.

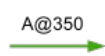
Applications of HLC: Database Consistency



Applications of HLC: Achieving strict consistency

- Consensus on the operation log and the written value: RAFT
- Concurrency control: contention issues with reader-writer locks
- Need for wait free transaction ordering
- Idea: Multi Version Concurrency Control (MVCC) with HLC timestamps

Operation: R(A, 350)



Key	Timestamp	Value
A	400	"current_value"
A	322	"old_value"
A	50	"original_value"
B	100	"value_of_b"


Result = "old_value"

Applications of HLC: Global Snapshots

- Request to obtain a snapshot of the data at a time t
- Idea of introducing dummy events between events e and f on the same node
- Every event has l between $l.e$ and $l.f$ with c as 0.
- Take snapshot at $l = t$ and $c = 0$.

Applications of HLC: Global Snapshots

- Consistent snapshot for $t = 10$



Applications of HLC: Event Ordering Service

- **Kronos:** EuroSys 2014 paper
- Can use HLC timestamps to schedule events and register with Kronos
- HLC properties ensure consistent cuts in the event trace
- Snapshots, Predicate detection etc. become local operations